

The Commons-Hub Model

A replicable model for organizing collaboratively-developed open source software around nonprofit and for-profit satellites

Status: Early concept draft — not reviewed by counsel. Nothing in this document should be treated as a legal or tax conclusion.

Scope: This document describes a generalized, clonable organizational pattern — it is written to apply to any group collaboratively developing and monetizing an open source commons, not to any single organization in particular.

Introduction: Why This, Why Now

Framing draft — the "Trust issues" thread below trails off in the original dictation and still needs to be finished; flagged rather than guessed at.

Corporations as a governance technology, not a law of nature

The corporation is not a naturally occurring phenomenon. It is one socially produced, historically specific answer to two recurring questions: when people organize together to produce something, who governs that production — and who benefits economically from it? Different historical eras have answered both questions differently: guilds, common land, [joint-stock charters](#), the industrial corporation, the modern platform company. Each is a governance structure for collective production and economic benefit, adopted — and later challenged — because the previous answer stopped sufficiently serving whoever held the power to change it. That has usually meant the elites of the era, not a broad public: the English enclosure movements discussed below were driven by landowners who benefited from privatizing common land, not by commoners demanding it; joint-stock charters emerged to mobilize capital for merchants and investors who needed a legal vehicle for it, not from popular pressure. The pattern worth naming plainly: governance structures tend to change when they stop working for whoever already has enough power to rewrite them.

This initiative proposes to treat that question as still open, and to design a next answer for it — specifically for a mode of production (collaboratively developed open source software) that the corporate form was never built to serve well in the first place, since corporate law defaults to concentrating both ownership and benefit rather than distributing them.

Why now: what AI changes, and what it doesn't

That question of who governs production and who benefits from it isn't just philosophical — economists have a specific answer for why the corporate form won out over more distributed alternatives in the first place, and it's worth being precise about it, because it's exactly what's shifting now. [Ronald Coase's 1937 answer](#) to "why do firms exist at all, if markets are efficient?" was **transaction costs**: coordinating complex production through a market — finding people, negotiating, verifying quality, for every task — is expensive, and firms exist because hierarchical management was, for most of industrial history, cheaper than market coordination for a given kind of production. Coase won the 1991 Nobel in Economics for this.

The clearest illustration is textile production just before the modern corporation took shape. Before Richard Arkwright's water frame (1769), English cloth was made under the *putting-out system* — pure market coordination: a merchant distributed raw wool or cotton to independent spinners and weavers working in their own homes, paid piece-rate for finished cloth, no employment relationship at all. It lost to the factory for three specific reasons, not a vague one: **production was unobservable** (a merchant only saw finished cloth, never the process, which produced chronic embezzlement of material and inconsistent quality); **scheduling was unenforceable** (dispersed workers set their own pace, often around farm work); and **the new machinery physically required centralization** — a water wheel could power a mill, not a cottage. Direct supervision solved the first two problems outright; centralization was the only way to use the new capital equipment at all.

AI doesn't erode all three of those reasons equally, and the honest version of this argument has to say so line by line rather than just asserting "AI makes things cheap":

- **Observability** — the strongest point of the parallel. AI-assisted code review, automated testing, and the trust/vetting mechanisms this document already builds in §6 do a real version of what a factory foreman did: make the production process legible without requiring everyone under one roof.
- **Capital-centralization — this one inverts, rather than just weakening.** In 1769 the machinery forced centralization; a spinner could never own a water wheel. Today the equivalent capital — compute, AI models — is rentable by the hour through cloud APIs. A dispersed contributor can access industrial-grade tooling *without* being inside a hierarchical firm that owns the capital equipment.
- **Scheduling/coordination** — largely already solved by tooling that predates this document (version control, CI/CD, issue trackers), with AI plausibly reducing the remaining friction further.
- **What doesn't get solved, and may get harder:** the 1770 embezzlement problem was about *material*; the equivalent risk now is trust in *contribution provenance* — can this code, and whoever submitted it, be trusted. Cheap AI-generated contribution volume doesn't shrink that problem, it

grows it — this is the [XZ Utils backdoor](#) risk from §6, restated as economic history rather than a security anecdote. **AI-driven abundance of code doesn't reduce the need for the trust layer in §6 — it increases it.**

One more honest qualifier, worth stating plainly rather than leaving implicit: "cheap" here means cheap in dollars and labor-hours, not cheap in resource terms. AI compute carries real energy, water, and embodied-carbon costs. An organization whose values already lean toward social and economic justice shouldn't present AI-driven abundance as costless — that would sit inconsistently with the rest of this document's posture, which flags rather than asserts wherever a claim is uncertain.

There's a second, independent reason hierarchy isn't needed here, distinct from anything AI changes. Coase's hierarchy doesn't only solve "I can't see what you're doing" — it solves "you don't actually want to do what I need," the **misalignment** between what a wage-earner wants (the largest wage for the least effort) and what a firm wants (the largest output for the wage), which is precisely what supervision hierarchies exist to manage. Self-selected, mission-driven contributors who opted in because they believe in the goal mostly don't have that gap to begin with — this isn't a cheaper way to supervise labor, it's a context where the thing supervision exists to fix is largely absent. This already has a name and a rigorous treatment, and it's a tighter fit here than Coase alone: [Yochai Benkler, "Coase's Penguin, or, Linux and The Nature of the Firm"](#) (Yale Law Journal, 2002) — the title is a direct response to Coase, using Linux as the empirical counter-case, and it argues for a third mode of production alongside markets and firms: **commons-based peer production**, where complex goods get built by self-organizing contributors motivated by non-monetary reasons (mission, reputation, craft), coordinated through shared infrastructure rather than either price signals or managerial hierarchy. This isn't speculative — it's the mode that already produced Linux and most of the open source stack this pattern itself depends on.

The coordination mechanism this points toward has a name too: [stigmergy](#) — agents coordinating indirectly through shared traces left in a common environment, rather than through direct communication or central command. An ant doesn't get orders; it reads a pheromone trail another ant left and reinforces or ignores it, and the colony converges on good paths without any ant ever computing one. It's a live concept in swarm robotics and multi-agent AI system design today, not just biology — and this document already specifies the pheromone trail: the shared Commons itself, the committer-standing record in §6, and the tactical voting weight in §7 are the shared environment that lets independent contributors coordinate toward a common goal without anyone directing traffic.


Which points to why dispersion is the better answer here, not merely an adequate one. Concentrating capability into one legal entity, one point of control, one target, is exactly the aircraft-carrier problem: enormous value in a single place is precisely what makes it worth attacking — legally (one entity to sue or deplatform), politically (one target to pressure), operationally (one point of funding or leadership failure). A dispersed, replicable structure doesn't have that failure mode — no single instance is load-

bearing, and if one node goes down, the pattern survives elsewhere. That is what the clonable, multi-Satellite shape in §1 is actually for: not just a way to scale, but a resilience strategy.

A return to the commons

Before the [enclosure movements](#) reshaped landholding in England, common land was governed collectively by the communities that worked it — not owned by any single lord, not managed by the state. Enclosure privatized that land, and with it, concentrated the benefit of centuries of collective stewardship into far fewer hands.

Garrett Hardin's later "[tragedy of the commons](#)" argument treated unowned shared resources as inherently doomed to overexploitation — but Elinor Ostrom's empirical research (which won her the [2009 Nobel Memorial Prize in Economic Sciences](#)) overturned that claim, documenting hundreds of real cases where communities successfully self-governed shared resources through their own institutional rules, without requiring either privatization or centralized state control.¹



Enclosure Act for Shifnal, 1793 (Shropshire Archives 539/1/5/3). Public domain, via Wikimedia Commons.

This initiative is, in effect, a proposal to apply Ostrom-style commons governance to a *digital* commons — open source software — rather than to land, water, or fisheries: a body of collectively produced software governed by the people who build it, with the Commons-Hub structure (§1 below) as the specific institutional design, and the funded-program-allocation [DAO](#) (§3–4) as one concrete rule within it, in Ostrom's sense of a community writing its own rules for how a shared resource's benefit gets distributed.

Enclosure Act for Shifnal, 1793 (Shropshire Archives 539/1/5/3). Public domain, via Wikimedia Commons.

Trust and bad-faith actors

Because the Commons is open source, nothing stops an organization or individual whose goals are directly hostile to the adopting organization's own mission from adopting the same tooling — no vetting mechanism changes that, since the license already grants everyone access regardless of standing in the community. What the organization can still control is narrower: who earns committer/maintainer standing on its own infrastructure, and who its community prioritizes when people ask for help. Both are handled informally, by ordinary maintainer judgment, rather than by a formal vetting system — see §6.

The present moment: elite overproduction and AI-driven displacement

Peter Turchin's [structural-demographic theory](#) identifies "elite overproduction" as a recurring precondition for social instability: when a society trains and credentials more aspirants for elite-track positions than it has positions to absorb them into, intra-elite competition intensifies, average outcomes for elite aspirants decline, and — critically — some fraction of those aspirants become "counter-elites," turning their training and ambition toward organizing opposition to the existing order rather than joining it.²



Peter Turchin, 2020. Photo: Peter Turchin, CC BY 4.0, via Wikimedia Commons.

The current U.S. software labor market is a live instance of this pattern: an education system that has spent two decades producing an increasing supply of highly credentialed software engineers, now colliding with AI-driven contraction of entry- and mid-level engineering hiring. A growing population of capable, credentialed, increasingly frustrated engineers — including recent graduates with advanced degrees — is finding the traditional elite-track path (a well-paid job at a major tech employer) narrowing or closing.

The opportunity

That population is not just a labor pool to hire from — it is a mobilizable base for a movement, in Turchin's counter-elite sense. This initiative's pitch to them is not purely economic: it offers a way to do real, technically substantial work, be paid for it through the mechanisms described in §3–4, and do so without having to take a job at an organization whose practices conflict with their politics. (Any organization adopting this pattern will generally already have its own public language naming the specific employers or practices it considers disqualifying — that language belongs in the adopting organization's own materials, not in this generalized document.) Positioned this way, the Commons-Hub pattern is simultaneously a funding/staffing strategy for whichever specific organization adopts it, and a replicable recruitment narrative for any future Satellite or Commons built on this pattern.


1. [Elinor Ostrom — Wikipedia](#)
2. [Elite overproduction — Wikipedia](#); [Structural-Demographic Theory — Peter Turchin](#)

1. Core Idea

At the center of the pattern is a **Commons**: a body of collaboratively developed open source software (and the engineering practices, standards, and shared libraries around it). The Commons is not itself a legal entity — it's an artifact and a body of practice. It functions as:

- an **engineering hub** — where core contributors maintain shared infrastructure, standards, and foundational libraries
- a **financial source and sync** — money flows *out* of the Commons (funding development of the shared stack) and *into* the Commons (royalties, IP licensing fees, or contribution-in-kind from satellites that build on it)

Any number of legal entities can be organized *around* a single Commons. A Commons might be a shared application framework, a set of foundational libraries, or a full product stack — the pattern is meant to generalize: any collaboratively-developed open source commons could sit at the center.



This multi-Satellite shape is a resilience property, not just a scaling one — see the Introduction ("Why now: what AI changes, and what it doesn't") for why dispersion is the better answer here, not merely an adequate one.

2. The Satellite Layer

Each **Satellite** is a [501\(c\)\(3\)](#) organized around some program of work that draws on the Commons. A Satellite may, as its revenue-generating programs mature, spin those programs out into a **wholly-owned for-profit subsidiary**: the nonprofit retains mission/governance control, while the subsidiary carries the operational/revenue-generating work.

Multiple Satellites can exist around one Commons simultaneously — e.g., one Satellite oriented toward labor organizing tools, another toward tenant organizing, another toward a specific geography — all built on the same shared stack, all feeding improvements and (where applicable) licensing revenue back to the Commons.

3. Governance Layer — Two Separate Mechanisms


It's important to keep these two things distinct; they are not the same mechanism and should not be conflated:

Organizational governance — the Board. Mission, priorities, protocols, and policy for each Satellite are set through ordinary nonprofit governance: the 501(c)(3)'s board of directors. **There is no [DAO](#) involved in setting priorities or policy.** The board decides what to build, which funded programs to pursue (grants, earned-revenue-funded initiatives, or internally allocated budget), and what the organization's direction is, exactly as any nonprofit board would.


Funded-program-allocation DAO — scoped narrowly, one mechanism among possibly several. A separate, purpose-built DAO mechanism governs exactly one thing: how the proceeds of an *already-secured* funded program — a grant, a commercial-revenue-funded initiative, or a board-allocated budget — are split among the individual developers (and pods) who contributed to earning it. Deliberately not scoped to grants alone: this organization does not intend to depend solely on grant funding long-term, and the mechanism doesn't care where the money came from. This DAO does not set policy, does not decide what to build, and does not govern the organization. Its entire remit is: given a funded program the board has already secured, and a defined pool of contributors who worked on it, determine what share of the disbursed funds each contributor or pod receives. Other narrowly-scoped DAO-like mechanisms may eventually exist for other specific purposes (to be defined later) — but priority-setting is never one of them.

4. Worked Example: Funded Program Allocation Mechanic

This is the concrete mechanic the allocation DAO implements. Note that every step here happens *after* the board has already decided to pursue and has secured the funding — the DAO has no role in that decision.



Part 1 — funding award through the vote.



Part 2 — the vote's output through payment.

Coordinape and Snapshot touch this pipeline only at the voting layer, and only lightly — both use a connected crypto wallet as a *signature/identity* mechanism, not as an on-chain computation. Snapshot votes are signed messages tallied off-chain (gasless); Coordinape's peer-allocation round is an ordinary web app. The genuine on-chain activity in this diagram is the treasury and payment layer (Safe multisig → Request Finance/Rise → developer wallet) — that was always going to be blockchain-based once the org chose to pay contributors in crypto, independent of the voting-tool choice.

1. A Satellite's board pursues and secures a funded program — e.g., a \$100,000 grant, an earned-revenue-funded initiative, or a board-allocated budget — on terms the board negotiated and remains accountable for.
2. Community members who want to contribute work toward the funded program **opt in** as contributors (no fixed roster; participation is self-selected by interest).
3. Each participating contributor completes standard tax intake (e.g., a [W-9](#) for US persons, [W-8BEN](#) for foreign contributors) and registers a wallet address for receiving funds (chain TBD — Bitcoin or

Ethereum are the current candidates). Intake must be complete *before* a contributor is eligible to receive any disbursement — this is a gate, not a follow-up step.

4. As the program hits its funded milestones, **active contributors vote** on how the released tranche is apportioned among those who did the work.
5. Votes can allocate a tranche either to an **individual contributor** or to a **pod** (a development sub-group working a piece of the program).
6. If a tranche is voted to a pod, the pod's own members then decide internally how to divide that chunk — the DAO-level vote only decides the pod's allocation, not the individual split within it.
7. The vote's output is an **allocation instruction**, not a payment. It feeds a disbursement pipeline that: confirms tax intake is on file for every recipient, executes payment only after that check passes, records each payment as compensation for services rendered on the specific funded program (not as a profit share or gift), tracks cumulative payments per contributor per calendar year, and issues [1099-NEC](#) (or the applicable equivalent) at year end for contributors who cross reporting thresholds. The DAO machinery is meant to be built integrated with this disbursement/tax pipeline from the start, not layered on top of it later.

5. Risks & Open Questions

Ranked roughly by how likely each is to force a redesign of the mechanic above.

LEGAL — MEDIUM SEVERITY

DAO scope must stay strictly allocation-only. Because the DAO governs only the *distribution* of already-secured funded-program dollars among contributors — not priorities, not policy, not whether to pursue the program — board fiduciary duty is much less implicated than a "DAO sets policy" design would be. But the board (or an authorized officer) should still retain audit/override authority over disbursements: the org, not the DAO, remains legally accountable for how those dollars are spent. Recommend the DAO's vote function as a binding allocation instruction that the org's disbursement process executes automatically, with the board retaining override/audit rights, not that the DAO's output be completely outside institutional control. **Worth confirming with counsel that this framing is sufficient**, but it is a materially smaller ask than earlier drafts of this document implied.

LEGAL / TAX — HIGH SEVERITY

Payments must be structured as compensation for services, not profit-sharing. Because contributors are being paid out of funded-program dollars administered by a 501(c)(3), each payment has to map to actual services that contributor performed on that program — this is the [private-inurement](#) guardrail. A peer vote determining *shares* is fine as a mechanism for sizing a services payment, but the underlying legal characterization must be "contractor payment for services performed" — documented as such (e.g., a lightweight statement of work or contribution record per contributor per milestone) — not "distribution of program winnings" or a profit-sharing distribution. **Requires confirmation from counsel.**

TAX

Intake must gate disbursement, not follow it. Building the DAO/wallet registration flow to require W-9/W-8 intake before funds are eligible to move keeps compliance built into the mechanism rather than bolted on after the fact — this is reflected as step 3 above and should stay a hard gate in implementation, not a best-effort follow-up.

TAX

Crypto valuation & withholding. Paying in crypto still requires fair-market-value USD conversion at time of payment for 1099 reporting purposes, and potentially backup withholding if a contributor fails to provide a W-9. Needs to be built into the disbursement pipeline, not handled ad hoc at year end.

LEGAL / TAX

Related-party transactions between Commons and Satellite. If the Commons is ever formalized as its own entity, money/IP flowing between it and a 501(c)(3) Satellite is a related-party transaction and needs arm's-length terms (e.g., an IP license at fair market value) to avoid private-benefit problems. Keeping the Commons deliberately entity-less (a shared codebase and set of practices, not a legal person) avoids this risk for as long as that's tenable.

LEGAL

DAO legal wrapper. Because the DAO's scope is now narrow (allocation only, with off-chain payment execution and board override), it may not need its own chartered legal vehicle (e.g., a [Wyoming DAO LLC](#)) the way a policy-setting DAO would — it could plausibly be implemented as an internal voting tool feeding an ordinary org-run disbursement process. Still **unclear/requires confirmation** whether any wrapper is needed, and if so, what liability exposure the voting mechanism itself carries.

OPERATIONAL

Chain/token choice. Bitcoin vs. Ethereum vs. something else — unresolved.

OPERATIONAL

Milestone verification. Who certifies that a funded milestone has actually been met, triggering a vote — likely a board or program-management function, but unresolved.

POLICY

Interaction with an adopting organization's existing comp philosophy. Many organizations publicly commit to transparent, formula-based compensation bands rather than manager/community discretion or negotiation. Reconciling that kind of stated philosophy with a peer-vote-based allocation mechanic — for anyone who is *also* a paid employee or contractor — needs to be resolved case by case per adopter.

SEQUENCING

Relationship to an interim fiscal sponsorship arrangement. This document describes a *permanent* structure. It is compatible with, and not blocked by, an adopting organization pursuing a fiscal sponsor to cover a gap while unincorporated — the sponsor relationship covers the near term;

this structure is a candidate for what the organization incorporates into (or graduates out of the sponsor into) later.

LEGAL — MEDIUM SEVERITY

Tactical governance tokens (§7) must be board-delegated authority, not personal property. A nonprofit board can lawfully delegate day-to-day/tactical decision-making to staff, a management team, or — per §7 — a token-weighted committer vote. What it cannot do is let a founder's personal, freely-held instrument operate as the org's real control mechanism independent of board authorization. The distinguishing factor is procedural: the board must formally adopt the token system as its own chosen delegation mechanism (a board resolution is enough) and retain the power to modify or unwind it — the same override/audit pattern already required of the funded-program-allocation DAO in Risk 1, one layer up. **Requires confirmation from counsel that the adoption resolution is sufficient**, particularly for the 501(c)(3) side; the for-profit subsidiary side is ordinary corporate practice and carries much less of this risk.

LEGAL — LOW SEVERITY

Securities-law exposure, walked through against Howey rather than just asserted. The four-part [Howey](#) test for an investment contract requires (1) an investment of money, (2) in a common enterprise, (3) a reasonable expectation of profit, (4) derived from the efforts of others. §7's tokens fail prong 1 (nothing is purchased — they're granted by a trusted peer, not bought), fail prong 3 (no profit-participation rights, and non-transferability means no resale value either — no dividends, no capital appreciation, nothing to expect a profit from), and prong 4 actively inverts (holders keep their weight only by remaining active contributors themselves — the opposite of the passive reliance on others' efforts Howey targets). Uniswap's UNI token — cited in an earlier draft of this risk item as a reason "governance only" doesn't automatically clear the bar — isn't the right comparison once transferability is off the table: UNI's exposure came from trading on open exchanges with a real, observable market price, which is exactly what non-transferability removes here. The direct precedent is [United Housing Foundation, Inc. v. Forman](#), 421 U.S. 837 (1975) — the Supreme Court held cooperative shares were *not* securities specifically because they couldn't be transferred outside the cooperative, weren't purchased for profit, and (notably, given §7's decay/vesting design) carried voting rights not tied to the raw number of shares held. Several states codify the same pattern by statute: California's [Consumer Cooperative Corporation Law](#) and [Limited Liability Worker Cooperative Act](#), and Illinois' [Limited Worker Cooperative Association](#)

[Act](#) (805 ILCS 317), both exempt non-transferable, participation-based cooperative membership interests from securities registration. This doesn't eliminate the question — **still requires confirmation from counsel**, per this document's own discipline of never treating a legal read as fully settled, and current SEC posture on digital-token no-action guidance specifically should be verified rather than assumed from older letters — but it's a genuinely low-probability finding, not an open question of comparable weight to Risk 2.

6. Community Trust and Contributor Standing

Because the Commons is open source, nothing stops a bad-faith actor from taking the code and building something the organization opposes — no vetting mechanism, cryptographic or otherwise, changes that; the license already grants everyone access regardless of standing in the community. (An earlier draft of this section specified a zero-knowledge/blockchain-based trust graph to address this. It was dropped: the mechanism couldn't actually prevent the harm it was aimed at — open-source code isn't excludable — and it added real engineering cost to solve a narrower problem than it looked like it was solving.)

What an organization *can* still control is much narrower, and doesn't need special machinery:

- **Committer/maintainer status on its own core infrastructure.** Handled the way most open-source projects already handle it: a new contributor earns commit access through sustained good contributions, vouched for by one or two people who already have it. Tracked however the maintainers already keep notes — a private list, a GitHub team, a line in a README — not a purpose-built trust protocol.
- **Community support prioritization.** In the org's community channels, **no one is ever denied help**. People the community already trusts get more effort, as a judgment call made in the moment by whoever's responding, not a lookup against a formal registry.

Both rely on ordinary human judgment and institutional memory, not on a system engineered to resist infiltration or graph harvesting — because there's no formal graph to infiltrate or harvest in the first place.

A secondary benefit: the support archive itself

Every request for help, whoever it comes from, either surfaces a real gap (a documentation hole, a bug) worth fixing regardless of the requester's intent, or adds to a growing record of previously-solved problems. That archive — mailing list threads, issue trackers, chat logs — is worth deliberately preserving and indexing on its own merits: it's a natural training/retrieval corpus for a self-serve support

tool (an LLM-backed FAQ bot, a searchable knowledge base) that reduces the org's support burden over time. This is a good idea independent of the trust question above, not a consolation prize for dropping the trust network.

7. Tactical Governance — Committer Voting Weight

This is a third, distinct mechanism, separate from both the board (§3, strategy/fiduciary oversight) and the funded-program-allocation DAO (§3–4, which splits the proceeds of an already-secured funded program). §7 governs a different question entirely: **day-to-day tactical decisions** — roadmap, technical direction, priorities — across both the nonprofit and the for-profit subsidiary. It carries no profit-participation rights; economic equity/compensation, if any, is a separate instrument handled through ordinary corporate mechanisms, not this one.

The mechanism

A founding steward starts by holding all tactical voting weight. It is distributed to developers who earn that steward's trust, and — this is the part that makes it scale rather than bottleneck through one person — those developers can further distribute weight to anyone else who is a registered committer in the ecosystem (§6). This reuses §6's existing eligibility gate rather than inventing a second identity system: to hold or receive tactical voting weight, you already have to have earned committer standing.

It cannot be bought or sold. No consideration changes hands in either direction — weight is granted based on trust, never purchased, and never transferable for payment. This is a deliberate design choice, not an afterthought: it's the single biggest factor in why this doesn't read as a security (see Risk 12), and it mirrors the closest real legal precedent for a non-transferable, participation-based governance interest — cooperative membership shares ([*United Housing Foundation v. Forman*](#), 1975; see Risk 12 for the full citation and analysis).

That restriction is enforceable only halfway by code, and the document shouldn't imply otherwise. A smart contract (or any access-control system) can check *who* is eligible to receive a transfer — is the recipient a registered committer — and reject transfers to anyone who isn't. It cannot see a side payment that happens off-chain: nothing stops someone from routing weight to a friend in exchange for cash or a favor arranged outside the system, and no code-level check can distinguish that from a good-faith grant of trust. Preventing that has to be a **norms-and-consequences problem** — a rule people are expected to follow, backstopped by discovery and revocation (of the tokens, and likely the offending parties' committer standing) when it's violated — not a technical guarantee the system provides on its own.

Two design constraints, both meant to prevent predictable failure modes

Freely transferable, permanently-held influence has a well-documented failure mode in every DAO/governance-token system that's tried it: weight accumulates in the hands of whoever collects the most of it, regardless of whether they're still contributing, and a bad early trust call compounds forever because there's no way to walk it back. Two constraints address this directly:

- **Vesting, not a lump sum.** A developer's allocation unlocks over time rather than landing all at once, so a bad early judgment call doesn't hand someone permanent outsized influence on day one.
- **Decay on inactivity.** Voting weight tracks *current* contribution, not historical accumulation. A committer who goes inactive for some period (window TBD — e.g., no merged contribution in a trailing six months) sees their weight decay toward zero over a following window, rather than retaining full influence indefinitely. Returning and contributing again should restore weight — exact re-affirmation mechanics are an open design parameter, not resolved here.

Also worth deciding, not yet resolved: a per-person concentration cap, so free redistribution can't be used to quietly assemble a controlling bloc through several small transfers.

The legal load-bearing point

For the 501(c)(3) side specifically, this system only works if it's characterized as **the board's own delegated, revocable authority** — not as the founder's personal property operating independently of the board. Concretely: the board formally adopts the token-weighted tactical-governance system as its chosen delegation mechanism (a board resolution), names an initial administrator, and retains the power to modify or unwind it. That's structurally identical to a board appointing and being able to remove an executive director — just implemented as a collective mechanism instead of a single appointed person. See Risks 11–12.

8. Economic Benefit — Sequencing Across Mechanisms

By this point there are three distinct mechanisms in play, and it's easy to conflate them since they all touch "who gets what." Worth seeing side by side once, rather than only encountering each in its own section:

Mechanism	Scope	Duration	Economic value	Who's eligible
DAO (§3–4)	Per funded program	Episodic — ends when the program does	Cash, paid for services rendered	Self-selected opt-in contributors

Mechanism	Scope	Duration	Economic value	Who's eligible
Tactical tokens (\$7)	Ongoing	Decays with inactivity	None — pure voice	Registered committers (\$6) who've earned trust
ESOP	Ongoing	Vests over years	Real equity	Legally must be broad-based — ~all full-time employees

A fourth candidate — an ESOP for the for-profit subsidiary — belongs on this table conceptually but not yet operationally. It only becomes worth evaluating once two things are both true, not on a calendar date:

1. **There's real enterprise value to distribute.** An ESOP requires an independent appraisal (no public market for the stock); with negligible or negative enterprise value, there's nothing meaningful to allocate — just administrative cost for its own sake.
2. **There's stable cash flow to fund the repurchase obligation.** Every vested share must eventually be bought back in cash when a participant leaves. A program-to-program or grant-to-grant cash position can't safely carry that liability; it takes predictable operating cash flow — in practice, real commercial revenue, not just grants.

Industry feasibility guidance (NCEO and others) puts typical setup costs at **\$100k–\$250k+**, with most transactions wanting **~\$1M+ in annual EBITDA and ~20+ employees** to justify that overhead — smaller ESOPs do exist, but this organization is well below even that floor today. One more sequencing note for later, not now: if a founder-exit tax deferral under [IRC §1042](#) is ever on the table, the subsidiary needs to already be a C-corp before that transaction — an entity-type decision worth making with this in mind well before it's actually needed.

The resulting ladder:

- **Now:** DAO for project-based cash payouts, tactical tokens for day-to-day voice. No ESOP — no enterprise value or stable cash flow yet to justify one, and no broad-based W-2 team to make "broad-based" mean anything.
 - **Once the subsidiary has sustained commercial revenue and real employees** (not just grant-funded opt-in contributors): ESOP feasibility becomes worth a real evaluation, running alongside — not replacing — the DAO and tactical tokens, each still doing its own job.
 - **If a founder-exit rollover is ever a goal:** C-corp status needs to already be in place before that transaction, which means the entity-type decision should account for this option early, not be revisited under time pressure later.
-

Draft prepared as a collaborative starting point. Next steps: legal/tax review of items 1–2 before any part of the funded-program allocation mechanic is implemented; legal review of items 11–12 before any part of the tactical governance token system (§7) is implemented.